

Distributed Machine Learning

Peacock

Yi Wang

Goal

- Scalability
 - engineering
 - mathematical

Gibbs Sampling

- Random initialization

words

text

Bill	Gates							
Bill	Gates	Microsoft						
		Microsoft	Windows					
				Steve	Jobs			
				Steve	Jobs	Apple		
						Apple	iPhone	
						Apple		iPad

Gibbs Sampling

- Update iteratively
 - the color of a word tend to be similar to other words in the same document.
 - the color of a word tend to be similar to its major color in the whole corpus.

Gibbs Sampling

- After convergence

words

text

Bill	Gates							
Bill	Gates	Microsoft						
		Microsoft	Windows					
				Steve	Jobs			
				Steve	Jobs	Apple		
						Apple	iPhone	
						Apple		iPad

Sufficient Statistics

- Random initialization

Bill	Gates							
Bill	Gates	Microsoft						
		Microsoft	Windows					
				Steve	Jobs			
				Steve	Jobs	Apple		
						Apple	iPhone	
						Apple		iPad

1	1
1	2
1	1
1	1
1	2
1	1
1	1

1	1	1		1	1	1		1
1	1	1	1	1	1	2	1	

Sufficient Statistics

- After convergence

Bill	Gates							
Bill	Gates	Microsoft						
		Microsoft	Windows					
				Steve	Jobs			
				Steve	Jobs	Apple		
						Apple	iPhone	
						Apple		iPad

2	
3	
2	
	2
	3
	2
	2

2	2	2	1					
				2	2	3	1	1

Many Documents

- Distribute by documents, duplicate model

Bill	Gates								2	
Bill	Gates	Microsoft							3	
		Microsoft	Windows						2	
				Steve	Jobs					2
				Steve	Jobs	Apple				3
						Apple	iPhone			2
						Apple		iPad		2

2	2	2	1					
				2	2	3	1	1

Many Documents

- Distribute by documents, duplicate model
 - n computers, each handle some documents and related statistics.
 - the model is duplicated on each computer.
 - computers sync up changes they made to local models.

Many Documents

- Sync-up local models
 - Synchronous can be done with MapReduce.
 - http://www.datalab.uci.edu/papers/distributed_topic_modeling.pdf
 - http://link.springer.com/chapter/10.1007/978-3-642-02158-9_26
 - Asynchronous sync-up needs self-made frameworks.
 - <http://papers.nips.cc/paper/3524-asynchronous-distributed-learning-of-topic-models>

Many Tokens

- Distribute by tokens.

Bill	Gates								2	
Bill	Gates	Microsoft							3	
		Microsoft	Windows						2	
				Steve	Jobs					2
				Steve	Jobs	Apple				3
						Apple	iPhone			2
						Apple		iPad		2
2	2	2	1							
				2	2	3	1	1		

Many Tokens

- Distribute by tokens, duplicate topic distributions.
- n computers, each handle some tokens, and maintains part of the model.
- the topic distributions is duplicated on each computer.
- computers sync up changes they made to topic distributions.

Many Documents and Tokens

- Distribute by documents and tokens.

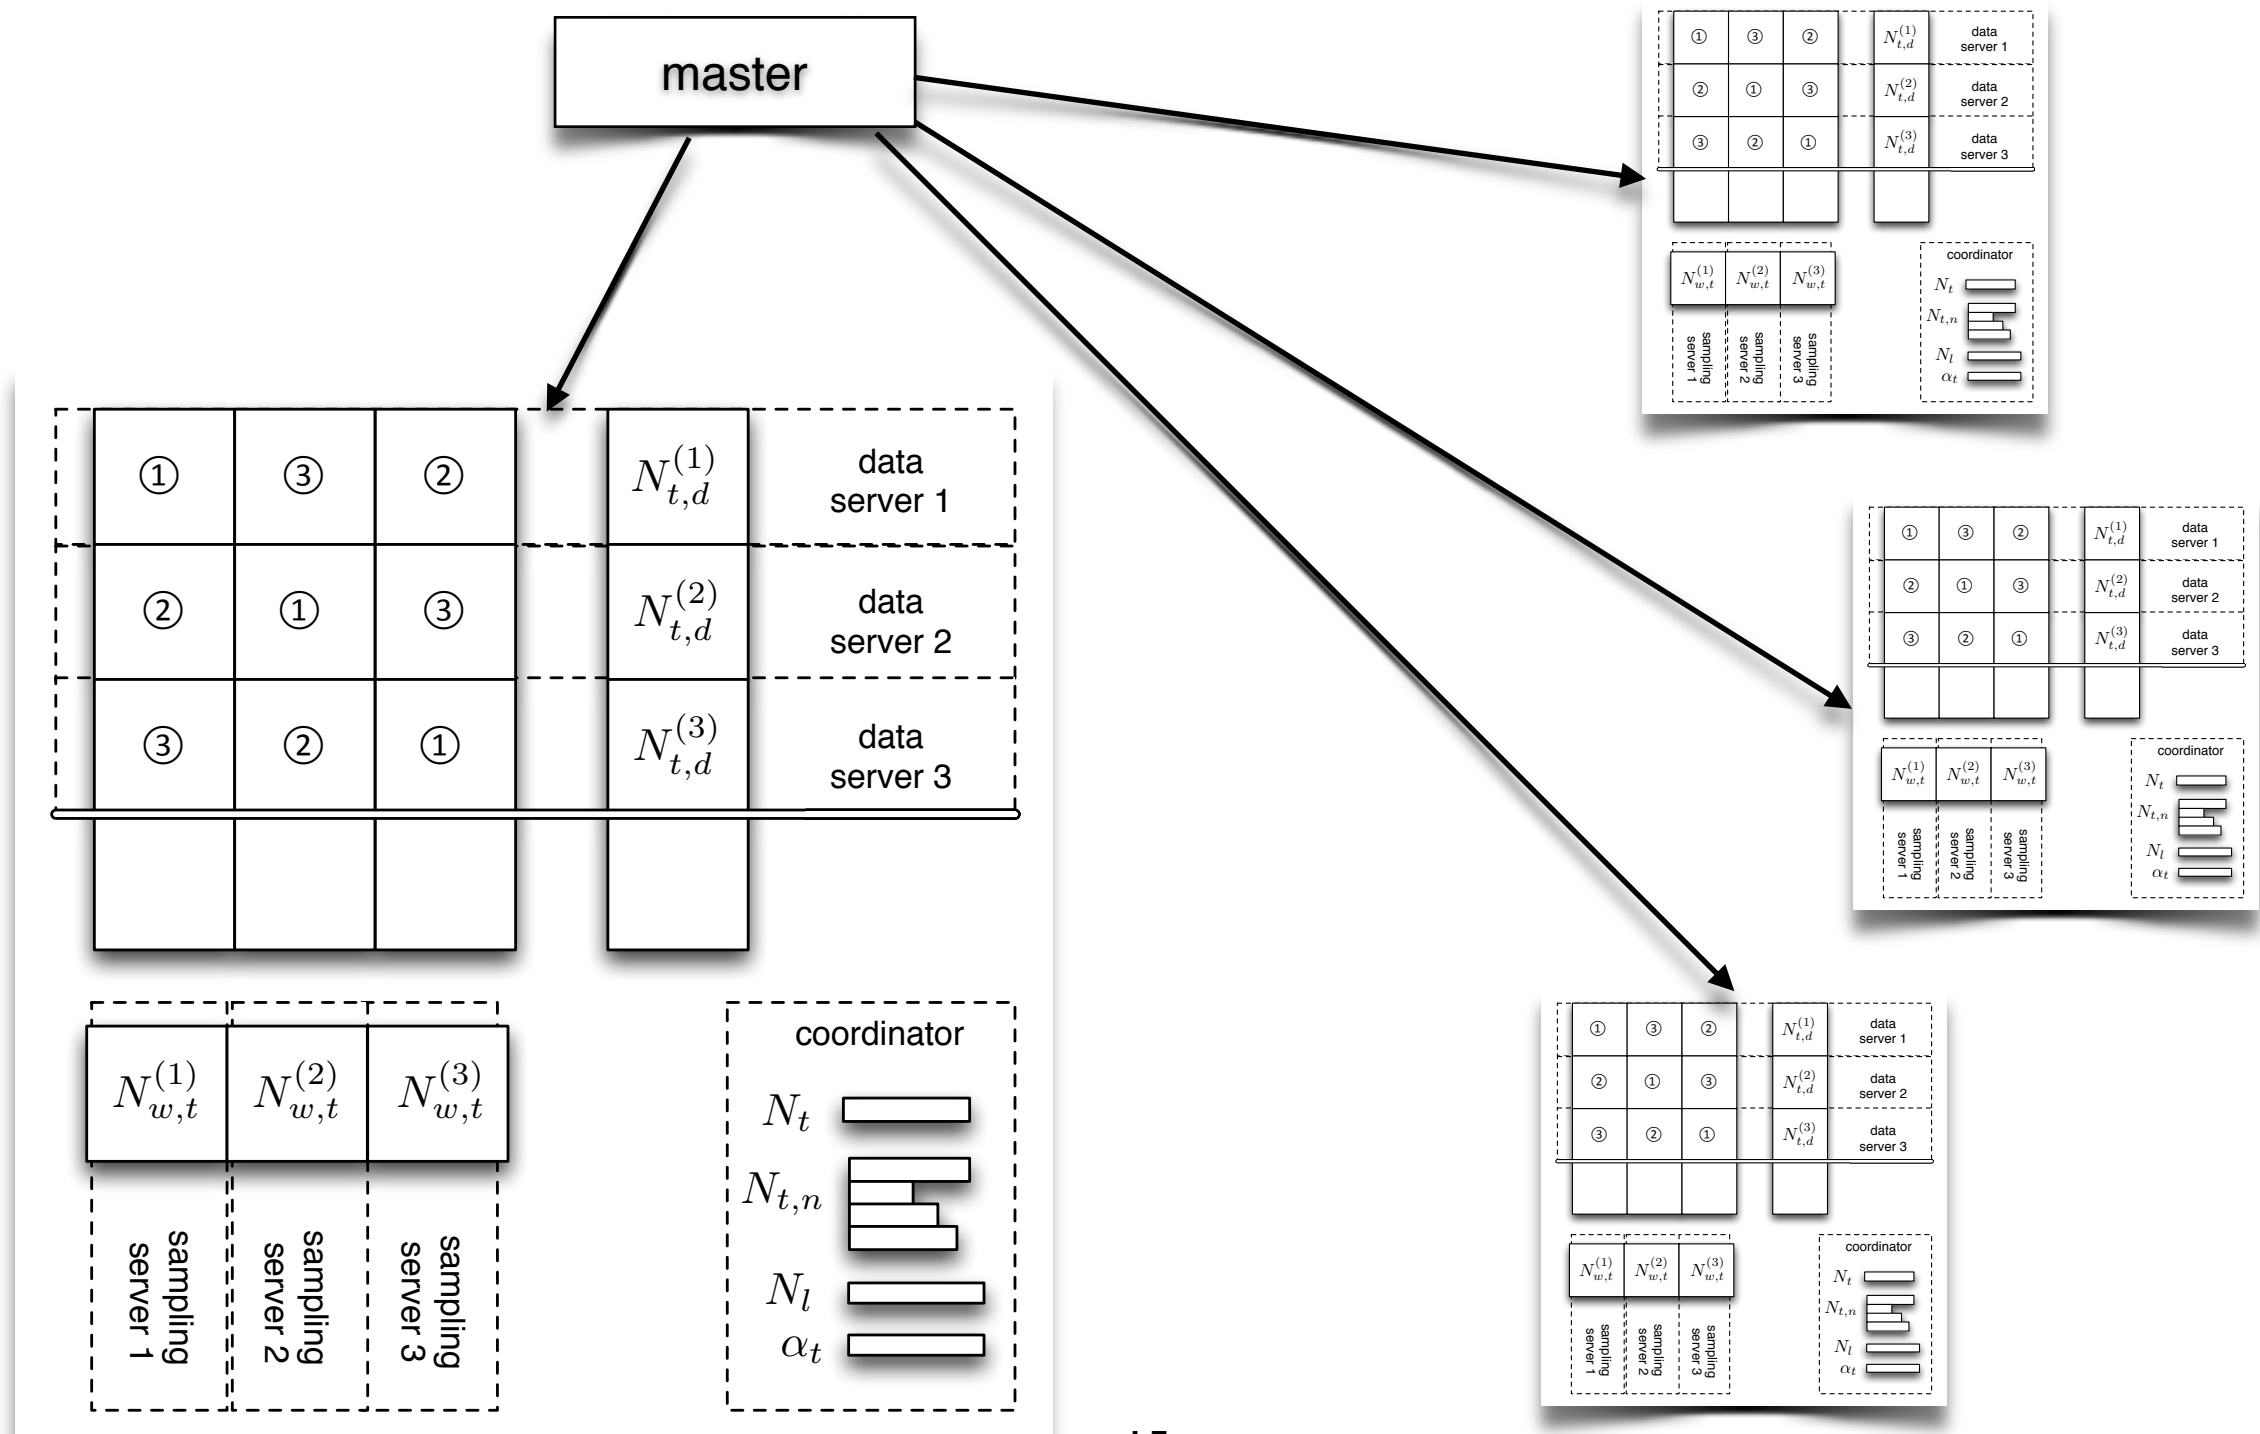
Bill	Gates								2	
Bill	Gates	Microsoft							3	
		Microsoft	Windows						2	
				Steve	Jobs					2
				Steve	Jobs	Apple				3
						Apple	iPhone			2
						Apple		iPad		2
2	2	2	1							
				2	2	3	1	1		

13

Many Documents and Tokens

- Distribute by documents and by tokens
 - n *loaders*, each handle some documents and related statistics.
 - another n *samplers*, each has a partition of the model.
 - An iteration has n steps.
 - In each step, a loader works with a sampler on a *block*.
 - **no model or distribution sync-up required.**

Scalable and Recoverable



Processes

- A *master* process, which controls all groups.
- m groups, each contains
 - a *coordinator* process,
 - n *loaders*, and
 - n *samplers*
- A group of n *aggregators*.

Master

- Ask Kubernetes to start m coordinators,
 - Maintain an *active* queue of M segments,
 - assign each coordinator a segment, and move assigned segments to *pending* queue.
- Waiting for coordinators' calls
 - If a coordinator finishes a segment, move it to *done*.
- watch these coordinators
 - if anyone died, restart it, assign pending segments.

Coordinator

- Ask Kubernetes to start n loaders and n samplers.
- Report to master and accept a segment (task):
 - for step $i = 0 \dots n-1$
 - loader x works with sampler $(x+i)\%n$ on updating block located at $x, (x+i)\%4n$.
 - each sampler report model updates to corresponding aggregator.
- If restarted, restart samplers and loaders, and samplers load model from aggregators.

Conclusion

- Done:
 - Modeling: asymmetric Dirichlet prior mimics Dirichlet process with huge K .
 - Engineering: asymmetric Dirichlet prior simplifies communication and sync-up so enables our architecture to learn huge K .
 - Automatically estimate K .
- Todo:
 - Extend Peacock to learn a deep hierarchical model.